

DSCC 3 & 4

UNIT - 3

Distributed File Systems

File Models

File Accessing, Sharing, and Caching

File Accessing

File Sharing

File Replication

Atomic Transactions

Case Study: HADOOP

Resource and Process Management

Task Assignment and Load Management

Task Assignment

Load Management

Unit - 4

UNIT-IV: Cloud Computing

Cloud Computing

Roots of Cloud Computing

Layers and Types of Clouds

Desired Features of a Cloud

Cloud Infrastructure Management

Infrastructure as a Service (IaaS)

Platform as a Service (PaaS)

Software as a Service (SaaS)

Challenges and Risks

Migrating into a Cloud

UNIT - 3

Here's a detailed breakdown with **10 points for each topic** in UNIT-III:

Distributed File Systems

1. Definition:

- A **Distributed File System (DFS)** enables file storage and access across a network of interconnected nodes, providing users with a unified file system interface.

2. Transparency:

- **Location Transparency:** Users don't need to know file location.
- **Replication Transparency:** Multiple replicas of files appear as a single entity.

3. **Scalability:**

- Designed to handle large datasets and an increasing number of users or nodes.

4. **Fault Tolerance:**

- Utilizes replication and recovery mechanisms to ensure data availability even during failures.

5. **Concurrency:**

- Multiple clients can access files simultaneously without interference.

6. **Performance:**

- Employs caching, replication, and optimized access strategies to reduce latency.

7. **Consistency:**

- Ensures all replicas of a file are updated uniformly after modifications.

8. **Access Control:**

- Provides user authentication and file permission mechanisms to secure data.

9. **Interoperability:**

- Supports multiple platforms and integrates with other file systems (e.g., NFS, SMB).

10. **Examples:**

- HDFS (Hadoop Distributed File System), Google File System (GFS), Amazon S3.

File Models

1. **Structured Model:**

- Files are organized in a predefined structure, such as tables or databases.

2. **Unstructured Model:**

- Files contain free-form data, such as text or multimedia, without predefined structure.

3. **Hierarchical Model:**

- Files are arranged in a tree-like directory structure for organization.

4. **Flat Model:**

- Files are organized in a single directory with no subfolders.

5. **Key Attributes:**

- File metadata includes name, size, creation time, and permissions.

6. **Access Types:**

- Read, write, append, and delete operations define how files are manipulated.

7. **Naming Conventions:**

- Uses unique identifiers or paths to locate files.

8. **Storage Locations:**

- Determines whether files are stored locally, on a server, or replicated across nodes.

9. **Access Granularity:**

- Byte-level or record-level access to data within a file.

10. **Examples:**

- Text files, binary files, log files, and structured XML/JSON files.
-

File Accessing, Sharing, and Caching

File Accessing

1. **Sequential Access:**

- Files are accessed in a linear sequence from start to end.

2. **Random Access:**

- Files can be accessed non-linearly using offsets or indices.

3. **Access Protocols:**

- Protocols like NFS, SMB ensure standardized file access in DFS.

4. **Remote Access:**

- Uses RPC (Remote Procedure Call) or APIs for accessing files stored on remote nodes.

5. **Caching Mechanisms:**

- Frequently accessed files are cached locally for quick access.

File Sharing

1. **Concurrent Access:**

- Multiple users or processes can access the same file simultaneously.

2. **Synchronization:**

- Locking mechanisms (e.g., read/write locks) to avoid conflicts.

3. **Version Control:**

- Keeps track of changes to shared files to maintain consistency.

4. **Access Control:**

- Implements permissions (read, write, execute) for shared files.

5. **Collaboration Tools:**

- DFS supports collaborative environments, e.g., Google Drive or Dropbox.
-

File Replication

1. **Definition:**

- Copies of the same file are maintained on multiple nodes.

2. **Purpose:**

- Improves fault tolerance and data availability.

3. **Replication Models:**

- **Synchronous:** Updates all replicas simultaneously.
- **Asynchronous:** Updates replicas at different times.

4. **Placement Policies:**

- Determines where replicas are stored to optimize performance and reliability.

5. **Consistency Models:**

- **Strong Consistency:** All replicas show the same data at any time.
- **Eventual Consistency:** Consistency achieved over time.

6. **Fault Tolerance:**

- Replication ensures data recovery in case of node failures.

7. **Load Balancing:**

- Distributes read and write requests across replicas to prevent overloading.

8. **Challenges:**

- Ensuring consistency, avoiding stale data, and managing overhead.

9. **Performance:**

- Balances storage cost with speed of access.

10. **Use Cases:**

- Data centers, HDFS, and RAID systems.
-

Atomic Transactions

1. **Definition:**

- Ensures that file operations are completed entirely or not at all.

2. **ACID Properties:**

- Atomicity, Consistency, Isolation, Durability.

3. **Role in DFS:**

- Maintains system integrity during concurrent file operations.

4. **Transaction Logs:**

- Maintains a log of operations to ensure rollback or recovery.

5. **Fault Tolerance:**

- Recovers incomplete transactions after crashes.

6. **Concurrency Control:**

- Prevents conflicts when multiple transactions modify the same data.

7. **Commit Protocols:**

- Ensures all participants in a distributed transaction agree on changes.

8. **Two-Phase Commit:**

- A protocol ensuring all nodes agree before committing a transaction.

9. **Rollback Mechanisms:**

- Reverts changes in case of transaction failure.

10. **Applications:**

- Banking systems, file replication in DFS.
-

Case Study: HADOOP

1. **Introduction:**

- Hadoop Distributed File System (HDFS) is a scalable, fault-tolerant DFS designed for big data.

2. **Architecture:**

- **NameNode:** Manages metadata.
- **DataNodes:** Stores file blocks.

3. **Fault Tolerance:**

- Replicates data across multiple nodes for reliability.

4. **Scalability:**

- Handles large datasets by distributing them over multiple nodes.

5. **Write-once-read-many:**

- Files are immutable once written, optimizing for reads.

6. **Data Blocks:**

- Files are divided into blocks (default: 128 MB) for distribution.

7. **Replication:**

- Default replication factor is 3 for fault tolerance.

8. **YARN:**

- Resource management and job scheduling framework in Hadoop.

9. **Use Cases:**

- Data warehousing, machine learning, and log analysis.

10. **Performance:**

- Suitable for high-throughput applications but not low-latency tasks.
-

Resource and Process Management

1. **Resource Allocation:**

- Ensures efficient usage of CPU, memory, and network bandwidth.

2. **Schedulers:**

- Allocates resources using FIFO, priority, or fair-share scheduling.

3. **Load Monitoring:**

- Tracks system resource usage to identify bottlenecks.

4. **Fault Recovery:**

- Reallocates failed tasks to healthy nodes.

5. **Process Migration:**

- Moves processes to underutilized nodes for better load distribution.

6. **Task Assignment:**

- Matches tasks to nodes based on resource availability.

7. **Job Prioritization:**

- Ensures high-priority jobs are executed first.

8. **Load Balancing:**

- Evenly distributes workload across nodes.

9. **Dynamic Allocation:**

- Adapts resource allocation in real-time based on workload changes.

10. **Tools:**

- Apache Mesos, Kubernetes for process and resource management.
-

Task Assignment and Load Management

Task Assignment

1. Static Assignment:

- Fixed allocation of tasks to nodes at design time.

2. Dynamic Assignment:

- Allocates tasks in real-time based on current system state.

3. Metrics Considered:

- CPU load, memory usage, network bandwidth.

4. Task Dependencies:

- Ensures dependent tasks are assigned to the same or nearby nodes.

5. Load Prediction:

- Uses historical data to forecast and allocate resources.

Load Management

1. Load Balancing:

- Redistributes workload to ensure no node is overburdened.

2. Load Sharing:

- Distributes excess load from busy nodes to idle nodes.

3. Algorithms:

- Round-robin, least connections, and weighted balancing.

4. Overhead Minimization:

- Reduces communication and migration costs during load balancing.

5. Real-time Adjustments:

- Continuously monitors and reallocates tasks as needed.

Let me know if you'd like further elaboration on any subtopic!

Unit - 4

UNIT-IV: Cloud Computing

Here's a detailed breakdown with **10 points for each topic** in UNIT-IV:

Cloud Computing

1. Definition:

- Delivery of computing resources (servers, storage, databases, networking, software) over the internet.

2. On-Demand Services:

- Provides resources dynamically without requiring upfront provisioning.

3. Resource Pooling:

- Combines resources for multiple users, enabling shared access.

4. Elasticity:

- Resources scale up or down based on demand.

5. Pay-as-You-Go:

- Users pay only for the resources they consume.

6. Virtualization:

- Key enabler of cloud computing; creates virtual instances of hardware and software.

7. Service Models:

- Includes IaaS, PaaS, and SaaS for varying levels of abstraction.

8. Deployment Models:

- Public, private, hybrid, and community clouds.

9. Security:

- Implements authentication, encryption, and compliance measures.

10. Examples:

- AWS, Microsoft Azure, Google Cloud Platform.
-

Roots of Cloud Computing

1. Utility Computing:

- Early model focused on providing computing as a metered utility.
 - 2. **Grid Computing:**
 - Combines distributed resources to solve complex problems.
 - 3. **Virtualization:**
 - Abstracts physical resources for efficient use in clouds.
 - 4. **Web Services:**
 - Enabled application integration and service-oriented architecture.
 - 5. **Data Centers:**
 - Centralized facilities hosting hardware and applications for cloud services.
 - 6. **Advancements in Networking:**
 - High-speed internet made cloud computing feasible.
 - 7. **Open Source Technologies:**
 - Projects like OpenStack and Kubernetes accelerated cloud adoption.
 - 8. **Economy of Scale:**
 - Reduced costs by centralizing and optimizing resources.
 - 9. **Internet Evolution:**
 - Transition from Web 1.0 to Web 2.0 supported cloud-hosted applications.
 - 10. **First Movers:**
 - Early pioneers like Salesforce (SaaS) and Amazon EC2 (IaaS).
-

Layers and Types of Clouds

1. **Layers:**
 - **Physical Layer:** Hardware infrastructure.
 - **Virtualization Layer:** Abstracts physical resources.
 - **Application Layer:** User-facing applications and interfaces.
2. **Types of Clouds:**

- **Public Cloud:** Open to the general public (e.g., AWS).
 - **Private Cloud:** Exclusive to one organization.
 - **Hybrid Cloud:** Combination of public and private.
 - **Community Cloud:** Shared by organizations with similar needs.
3. **Infrastructure Layer:**
 - Includes servers, storage, and networking.
 4. **Platform Layer:**
 - Provides tools and frameworks for application development.
 5. **Application Layer:**
 - Hosts software and services for end-users.
 6. **Service Models:**
 - IaaS, PaaS, SaaS correspond to infrastructure, platforms, and software layers.
 7. **Deployment Differences:**
 - Public clouds are cost-effective; private clouds offer more control.
 8. **Virtualization Technologies:**
 - Hypervisors like VMware, KVM, and Xen underpin cloud layers.
 9. **Management Tools:**
 - Includes tools like Terraform and Kubernetes.
 10. **Scalability and Flexibility:**
 - Different cloud types support various business needs.
-

Desired Features of a Cloud

1. **Scalability:**
 - Dynamic scaling of resources based on demand.
2. **Reliability:**
 - Redundant systems ensure high availability.
3. **Security:**

- Data encryption, firewalls, and compliance measures.
 - 4. **Cost Efficiency:**
 - Pay-as-you-go pricing reduces upfront costs.
 - 5. **Accessibility:**
 - Accessible from any device with an internet connection.
 - 6. **Automation:**
 - Provisioning and scaling handled automatically.
 - 7. **Multi-tenancy:**
 - Serves multiple users securely on shared infrastructure.
 - 8. **Interoperability:**
 - Supports integration with other systems and services.
 - 9. **Performance:**
 - High-speed processing and low-latency networking.
 - 10. **Self-Service:**
 - Users can provision and manage resources independently.
-

Cloud Infrastructure Management

1. **Definition:**
 - Managing physical and virtual resources in a cloud environment.
2. **Monitoring:**
 - Tracks system performance and usage.
3. **Provisioning:**
 - Allocates resources for specific tasks or users.
4. **Configuration Management:**
 - Automates setup and maintenance of infrastructure.
5. **Resource Optimization:**
 - Ensures efficient usage of CPU, memory, and storage.
6. **Security Management:**

- Enforces access control, encryption, and threat monitoring.

7. Backup and Recovery:

- Ensures data integrity and availability during failures.

8. Compliance:

- Meets regulatory requirements for data handling.

9. Automation Tools:

- Tools like Ansible, Puppet, and Chef simplify management.

10. Cost Management:

- Monitors resource usage to optimize expenses.
-

Infrastructure as a Service (IaaS)

1. Definition:

- Provides virtualized computing resources over the internet.

2. Features:

- Compute, storage, and networking on-demand.

3. Use Cases:

- Hosting websites, running enterprise applications.

4. Examples:

- AWS EC2, Google Compute Engine.

5. Elasticity:

- Easily scale resources based on workload.

6. Customization:

- Full control over operating systems and applications.

7. Security:

- Responsibility shared between provider and user.

8. Cost Model:

- Pay-per-use; reduces CapEx.

9. Benefits:

- Flexible, scalable, and fast deployment.

10. **Challenges:**

- Requires expertise to manage infrastructure effectively.
-

Platform as a Service (PaaS)

1. **Definition:**

- Provides platforms for developers to build, test, and deploy applications.

2. **Features:**

- Includes frameworks, databases, and development tools.

3. **Examples:**

- Google App Engine, Microsoft Azure App Services.

4. **Ease of Use:**

- Simplifies development by abstracting infrastructure management.

5. **Scalability:**

- Automatically scales resources with application demand.

6. **Focus:**

- Developers focus on code; providers handle infrastructure.

7. **Customization:**

- Limited compared to IaaS but sufficient for most apps.

8. **Security:**

- Applications inherit provider's robust security measures.

9. **Cost Model:**

- Pay-as-you-use for services consumed.

10. **Challenges:**

- Vendor lock-in due to proprietary tools.
-

Software as a Service (SaaS)

1. **Definition:**

- Delivers software over the internet; users access via web browsers.

2. Features:

- No installation or maintenance required by users.

3. Examples:

- Google Workspace, Salesforce, Dropbox.

4. Accessibility:

- Accessible from any device with a browser.

5. Cost Model:

- Subscription-based pricing.

6. Multi-tenancy:

- Single instance serves multiple users.

7. Updates:

- Automatically updated by the provider.

8. Security:

- Provider ensures compliance and threat protection.

9. Use Cases:

- CRM, email, collaboration tools.

10. Challenges:

- Data security concerns and vendor dependency.
-

Challenges and Risks

1. Security Concerns:

- Data breaches, compliance, and insider threats.

2. Downtime:

- Service outages can disrupt business operations.

3. Vendor Lock-In:

- Limited ability to switch providers without re-architecting.

4. Cost Overruns:

- Poor management of resources leads to unexpected expenses.

5. Data Privacy:

- Risk of unauthorized access to sensitive data.

6. Integration Issues:

- Difficulties in connecting with on-premise systems.

7. Performance:

- Latency issues due to network dependency.

8. Scalability Limits:

- Some providers may not meet peak demands effectively.

9. Legal and Compliance:

- Challenges adhering to international data laws.

10. Complexity:

- Managing multi-cloud environments increases operational complexity.
-

Migrating into a Cloud

1. Definition:

- Transitioning from traditional IT infrastructure to cloud-based services.

2. Benefits:

- Reduces costs, improves scalability, and enables innovation.

3. Planning:

- Requires careful assessment of workloads and dependencies.

4. Broad Approaches:

- **Rehosting:** Lift-and-shift existing applications.
- **Refactoring:** Optimize applications for cloud environments.

5. Seven-Step Model:

- Assess, choose provider,

plan migration, secure data, execute migration, test, optimize.

1. Security:

- Encryption and access controls during migration.

2. **Tools:**

- Use tools like AWS Migration Hub for streamlined transitions.

3. **Challenges:**

- Downtime, data loss risks, and compatibility issues.

4. **Hybrid Approach:**

- Retain critical workloads on-premise while migrating others.

5. **Testing:**

- Post-migration testing ensures everything functions as expected.