

Short Notes DSCC

UNIT-I

- ☐ Introduction to Distributed Systems:
 - ☐ Characteristics of Distributed Systems-Introduction,
 - ☐ Examples of Distributed systems (Client server,
 - ☐ peer to peer,
 - ☐ grid and cloud computing),
 - ☐ Advantages of distributed systems,
 - ☐ System models - Introduction,
 - ☐ Architectural and Fundamental models,
 - ☐ Networking and Internetworking,
 - ☐ Interposes Communication (message passing and shared memory),
 - ☐ Distributed objects and Remote Method Invocation,
 - ☐ RPC,
 - ☐ Events and notifications,
 - ☐ Case study-Java RMI.
 - ☐ UNIT-II
 - ☐ Synchronization:
 - ☐ Time and Global States-Introduction,
 - ☐ Logical Clocks,
 - ☐ Synchronizing physical clocks events and Process states,
 - ☐ Synchronizing physical clocks,
 - ☐ logical time and logical clocks,
 - ☐ global states,
 - ☐ distributed debugging,
 - ☐ Coordination and Agreement:
 - ☐ Distributed mutual exclusion,
 - ☐ Elections,
 - ☐ Multicast communication,
 - ☐ consensus and related problems.
 - ☐ Here are comprehensive notes with **10 points** for each topic in the **Distributed Systems and Cloud Computing (DSCC)** syllabus.
-

UNIT I: Introduction to Distributed Systems

1. Characteristics of Distributed Systems:

1. **Concurrency:** Supports multiple simultaneous processes and user interactions.
2. **No Global Clock:** Lacks a unified time reference, requiring synchronization mechanisms.
3. **Independent Failures:** Individual node failures do not cause the entire system to fail.
4. **Scalability:** Can efficiently grow by adding more nodes to handle increased loads.
5. **Heterogeneity:** Integrates diverse hardware and software platforms seamlessly.
6. **Transparency:** Users experience the system as a single entity, unaware of the underlying complexity.
7. **Resource Sharing:** Enables the sharing of resources (e.g., files, processing power) among nodes.
8. **Openness:** The system is open to new components and services, allowing for integration and flexibility.
9. **Security:** Requires mechanisms for secure communication and data protection across nodes.
10. **Interactivity:** Allows real-time interaction between users and services, enhancing user experience.

2. Examples of Distributed Systems:

1. **Client-Server Model:** Servers provide resources or services to multiple clients (e.g., web services).
2. **Peer-to-Peer (P2P):** All nodes share equal responsibility and resources (e.g., file-sharing networks like BitTorrent).
3. **Grid Computing:** Combines resources from various locations to perform large-scale computations (e.g., SETI@home).
4. **Cloud Computing:** Delivers on-demand resources and services over the internet (e.g., AWS, Google Cloud).
5. **Distributed Databases:** Multiple databases located at different sites but provide a unified view (e.g., NoSQL databases).
6. **Microservices Architecture:** Applications are built as a collection of loosely coupled services (e.g., Netflix architecture).
7. **Distributed File Systems:** File systems that allow files to be stored across multiple servers (e.g., Google File System).
8. **Sensor Networks:** Networks of distributed sensors for monitoring environmental data (e.g., smart cities).
9. **Blockchain:** A distributed ledger technology that allows multiple parties to maintain a shared database (e.g., Bitcoin).
10. **Content Delivery Networks (CDN):** Distributes content across multiple servers to improve access speed (e.g., Akamai).

3. Advantages of Distributed Systems:

1. **Improved Performance:** Increased throughput and response times due to parallel processing.
2. **Fault Tolerance:** Redundancy and replication ensure system reliability in case of failures.
3. **Scalability:** Easy to add resources without downtime, allowing growth with demand.
4. **Resource Utilization:** Efficient use of resources by pooling together distributed nodes.
5. **Cost Efficiency:** Leverages commodity hardware to reduce infrastructure costs.
6. **Geographical Distribution:** Resources can be located near users to decrease latency.
7. **Flexibility:** Can accommodate different services and components dynamically.
8. **Ease of Maintenance:** Individual components can be updated or replaced without system-wide disruptions.

9. **Enhanced Collaboration:** Facilitates collaboration across organizations or departments by sharing resources.
10. **Modularity:** Systems can be designed with modular components, making it easier to develop and maintain.

4. System Models:

1. **Architectural Models:**
 - Client-Server
 - Peer-to-Peer
 - Multi-tier Architecture
2. **Fundamental Models:**
 - Interaction Models: Describes how processes communicate (synchronous/asynchronous).
 - Failure Models: Types of failures like crash, omission, and timing failures.
3. **Communication Models:** Defines how processes share information (message passing vs. shared memory).
4. **Consistency Models:** Determines how updates are propagated and viewed (strong vs. eventual consistency).
5. **Security Models:** Addresses authentication, authorization, and data integrity.
6. **Replication Models:** Strategies for replicating data across nodes for availability.
7. **Location Models:** Determines how nodes and resources are located and accessed.
8. **Resource Management Models:** How resources are allocated and managed among processes.
9. **Data Distribution Models:** Methods for distributing data across nodes (partitioning and replication).
10. **Service-Oriented Architecture (SOA):** A design approach based on the use of services to support software.

5. Networking and Internetworking

1. **Network Topologies:** Describes how nodes are connected (star, ring, mesh).
2. **Protocols:** Defines rules for communication (TCP/IP, HTTP, FTP).
3. **Network Types:** Distinguishes between LAN, WAN, MAN, and the internet.
4. **Routing and Switching:** Techniques for directing data across networks.
5. **Network Security:** Protects data in transit through encryption and secure protocols (SSL/TLS).
6. **Quality of Service (QoS):** Ensures reliable service and performance metrics.
7. **Bandwidth Management:** Controls data flow to prevent congestion and ensure efficient use of resources.
8. **Network Addressing:** IP addressing schemes and the role of DNS.
9. **Network Management:** Tools and techniques for monitoring and managing network performance.
10. **Interoperability:** Ability of different systems to work together across various networks.

6. Interprocess Communication:

1. **Message Passing:** Processes communicate by sending and receiving messages, supporting both synchronous and asynchronous modes.
2. **Shared Memory:** Processes share a common memory space for communication, requiring synchronization mechanisms (e.g., semaphores).
3. **Remote Procedure Calls (RPC):** Allows executing procedures on remote machines as if they were local calls.
4. **Sockets:** Provides an interface for network communication between processes over TCP/IP.
5. **Streams:** Continuous flow of data between processes, supporting byte-oriented communication.
6. **Pipes:** Unidirectional data channels for communication between processes, often used in UNIX systems.
7. **Events and Notifications:** Mechanisms to notify processes about changes or occurrences (e.g., publish/subscribe models).
8. **Signals:** Notifications sent to processes to indicate events (e.g., interrupt handling).
9. **Inter-Thread Communication:** Mechanisms like mutexes and condition variables for communication between threads within the same process.
10. **Data Serialization:** Converting data structures into a format suitable for transmission (e.g., JSON, XML).

7. Distributed Objects and Remote Method Invocation (RMI)

1. **Distributed Objects:** Objects that can be accessed from multiple nodes in a distributed system.
2. **Remote Method Invocation (RMI):** Java-based technology that allows method calls on remote objects as if they were local.
3. **Stubs and Skeletons:** Client-side and server-side components in RMI that facilitate communication.
4. **Marshalling and Unmarshalling:** Process of packaging data for transmission and extracting it at the destination.
5. **Java Naming and Directory Interface (JNDI):** Used to look up remote objects in RMI.
6. **Exception Handling:** Managing exceptions that occur during remote method calls.
7. **Security Policies:** Configuring security settings for RMI applications (e.g., access control).
8. **Serialization:** Converting objects into a byte stream for transmission over a network.
9. **RMI Registry:** A service that allows remote objects to be registered and looked up by clients.
10. **Performance Considerations:** Issues related to latency and throughput in RMI.

8. RPC (Remote Procedure Call):

1. **Definition:** A protocol that allows a program to execute a procedure on a remote server as if it were local.
2. **Client-Server Architecture:** RPC involves a client sending a request to a server to execute a function.
3. **Stub Mechanism:** Client-side stub prepares the request, while server-side stub processes it.
4. **Marshalling:** Packing the parameters into a message format for transmission.
5. **Unmarshalling:** Extracting the parameters on the server side.
6. **Transport Protocols:** Commonly uses TCP or UDP for communication.
7. **Error Handling:** Mechanisms to handle errors in remote calls (e.g., timeouts, retries).
8. **Synchronous vs. Asynchronous RPC:** Synchronous blocks the client until the response is received; asynchronous allows the client to continue processing.
9. **Security:** Authentication and encryption are vital for securing RPC communications.
10. **Performance Issues:** Network latency and overhead can affect the efficiency of RPC.

9. Events and Notifications:

1. **Event-Driven Architecture:** Components communicate by producing and responding to events.
2. **Event Producers and Consumers:** Producers generate events, and consumers act upon them.
3. **Publish-Subscribe Model:** Decouples event producers from consumers, allowing for flexible communication.
4. **Notification Services:** Systems that manage event notifications and delivery (e.g., JMS).
5. **Event Queues:** Buffering events for processing by consumers.
6. **Asynchronous Communication:** Events are processed independently of the producer's execution flow.
7. **Event Filtering:** Mechanisms to control which events are sent to which consumers.

8. **Event Processing:** Techniques for analyzing and responding to events in real-time (e.g., Complex Event Processing).
9. **Reliability and Durability:** Ensuring events are not lost during transmission or processing.
10. **Use Cases:** Common in systems like messaging apps, real-time monitoring, and

notification systems.

10. Case Study – Java RMI

1. **Overview of Java RMI:** Allows Java applications to invoke methods on objects located remotely.
 2. **Architecture:** Involves client stubs, server skeletons, and RMI registry for object discovery.
 3. **Use Cases:** Ideal for building distributed applications that require remote object interaction.
 4. **Serialization:** Automatic handling of object serialization for parameter passing.
 5. **Exception Handling:** Provides mechanisms for handling remote exceptions.
 6. **Security Features:** Supports SSL and access control for secure remote communications.
 7. **JNDI Integration:** Allows looking up remote objects by name.
 8. **Configuration:** Setup of RMI services, including port configuration and security policies.
 9. **Performance:** Considerations for optimizing remote calls (e.g., caching).
 10. **Limitations:** Constraints related to network latency and complexity of managing distributed objects.
-

UNIT II: Synchronization

1. Time and Global States:

1. **Global State:** Represents the state of a distributed system at a given time.
2. **Snapshot Algorithm:** Captures a consistent global state of the system without halting processes.
3. **Causality:** Understanding the cause-effect relationship between events in a distributed system.
4. **Logical Time:** Uses a counter to order events without relying on physical clocks (e.g., Lamport timestamps).
5. **Physical Time:** Actual wall-clock time used to coordinate events across distributed nodes.
6. **Global Clock Synchronization:** Techniques to synchronize physical clocks across distributed nodes (e.g., NTP).
7. **Logical Clocks:** Mechanisms to order events based on logical time (e.g., vector clocks).
8. **Distributed Debugging:** Techniques for debugging distributed systems, often using global states and event tracing.
9. **State Consistency:** Ensuring that global states reflect a consistent view across all nodes.
10. **Time-Stamping:** Assigning timestamps to events for ordering and consistency.

2. Logical Clocks:

1. **Definition:** A mechanism to order events in a distributed system without a global clock.
2. **Lamport Timestamps:** Assigns a numerical timestamp to each event to maintain a causal ordering.
3. **Vector Clocks:** An extension of Lamport timestamps to capture causality among distributed events.
4. **Clock Synchronization:** Techniques to synchronize logical clocks across nodes.
5. **Ordering Events:** Determines the causal relationship between events based on timestamps.
6. **Event Categories:** Classifies events into concurrent, causal, or unrelated based on timestamps.
7. **Algorithm Implementation:** Common algorithms for implementing logical clocks.
8. **Applications:** Used in distributed databases, version control systems, and concurrent programming.
9. **Limitations:** Challenges related to overhead and complexity in maintaining logical clocks.
10. **Comparison:** Differences between logical clocks and physical clocks in distributed systems.

3. Synchronizing Physical Clocks:

1. **Clock Drift:** The gradual deviation of a clock from the correct time due to inaccuracies.
2. **Network Time Protocol (NTP):** A widely used protocol for synchronizing clocks over a network.
3. **Time Servers:** Special servers that provide time information to synchronize client clocks.
4. **Stratum Levels:** NTP hierarchy indicating the distance from an authoritative time source.
5. **Synchronization Algorithms:** Techniques for adjusting local clocks based on server time.
6. **Precision and Accuracy:** Measures of how closely synchronized clocks are to the actual time.
7. **Latency Considerations:** Network delays that affect synchronization accuracy.
8. **Clock Adjustment Techniques:** Methods to correct clock differences (e.g., stepping, slewing).
9. **Synchronization Challenges:** Issues such as network delays, variable latency, and failures.
10. **Applications:** Importance of synchronized clocks in financial systems, telecommunications, and distributed databases.

4. Logical Time and Logical Clocks

1. **Logical Time:** A way to order events based on their occurrence rather than actual time.
2. **Event Ordering:** Establishes a sequence for events that is independent of physical time.
3. **Lamport Timestamps:** A simple method of implementing logical time using counters.
4. **Vector Clocks:** Extends Lamport timestamps to capture more complex causal relationships.
5. **Concurrency Detection:** Identifying concurrent events based on logical clock values.
6. **Causality Tracking:** Determining which events influence others using logical clocks.
7. **Implementation Challenges:** Complexity in maintaining logical clocks across distributed nodes.
8. **Impact on Consistency:** Logical clocks contribute to ensuring consistency in distributed systems.
9. **Applications:** Used in distributed databases, messaging systems, and collaborative applications.
10. **Comparison:** Logical clocks vs. physical clocks in maintaining order and causality.

5. Global States:

1. **Definition:** The complete state of a distributed system at a specific point in time.
2. **Consistency Models:** Rules for maintaining a consistent global state across distributed nodes.
3. **Snapshot Algorithms:** Techniques for capturing global states without stopping processes.
4. **Causal Relationships:** Understanding how events affect the global state.

5. **State Capture:** Methods for collecting state information from all nodes.
6. **Global Snapshot Properties:** Must reflect a consistent view of the system.
7. **Applications:** Used in debugging, fault tolerance, and state recovery.
8. **Challenges:** Difficulty in capturing a consistent global state due to concurrent processes.
9. **State Representation:** How global states are represented and stored.
10. **Impact on System Design:** Considerations for designing systems to maintain global state consistency.

6. Distributed Debugging:

1. **Definition:** Techniques for debugging programs running on distributed systems.
2. **Challenges:** Complexity due to concurrency, communication delays, and lack of global state visibility.
3. **Event Tracing:** Recording events to analyze system behavior and interactions.
4. **Global States in Debugging:** Using snapshots to understand system state at a specific time.
5. **Causal Analysis:** Understanding how events relate to one another in distributed debugging.
6. **Debugging Tools:** Software and frameworks that assist in debugging distributed applications.
7. **Monitoring and Logging:** Collecting data from distributed processes for analysis.
8. **Replay Mechanisms:** Techniques for replaying events to understand issues.
9. **Error Localization:** Identifying the source of errors in distributed systems.
10. **User Interfaces:** Tools for visualizing and interacting with distributed debugging data.

7. Coordination and Agreement

1. **Definition:** Mechanisms that allow distributed processes to coordinate their actions and reach agreements.
2. **Distributed Mutual Exclusion:** Techniques to ensure only one process can access a resource at a time.
3. **Election Algorithms:** Processes for electing a coordinator in distributed systems (e.g., Bully algorithm).
4. **Consensus Protocols:** Algorithms for reaching agreement among distributed processes (e.g., Paxos, Raft).
5. **Multicast Communication:** Sending messages to multiple recipients simultaneously.
6. **Atomic Broadcast:** Guarantees that messages are delivered to all processes in the same order.
7. **Resource Allocation:** Strategies for allocating resources in a distributed environment.
8. **Failure Handling:** Mechanisms for managing process failures during coordination.
9. **Consistency Models:** Ensuring all processes agree on the same state or value.
10. **Use Cases:** Applications of coordination and agreement in distributed databases, cloud computing, and collaborative systems.

8. Distributed Mutual Exclusion:

1. **Definition:** A mechanism to ensure that multiple processes do not access a critical section simultaneously.
2. **Centralized Approach:** A central coordinator grants access to the critical section (e.g., centralized locking).
3. **Token Ring Algorithm:** A token circulates in a logical ring, granting access to the critical section.
4. **Ricart-Agrawala Algorithm:** A request-based approach using timestamps for granting access.
5. **Lamport's Bakery Algorithm:** A ticket-based mechanism for mutual exclusion.
6. **Performance Metrics:** Evaluating efficiency based on response time and resource utilization.
7. **Scalability:** Ability to maintain mutual exclusion with increasing numbers of processes.
8. **Fairness:** Ensuring all processes get a chance to access the critical section.
9. **Deadlock Prevention:** Mechanisms to avoid deadlocks in mutual exclusion protocols.
10. **Use Cases:** Important in databases, file systems, and concurrent programming.

9. Elections:

1. **Definition:** Algorithms to select a leader or coordinator in distributed systems.
2. **Bully Algorithm:** Processes with higher IDs claim leadership, forcing lower IDs to step down.
3. **Ring Algorithm:** A token circulates in a logical ring, and processes vote to select a leader.
4. **Performance Metrics:** Efficiency and speed of the election process.
5. **Failure Handling:** Detecting and handling failures of the elected leader.
6. **Fairness:** Ensuring all processes have an equal opportunity to become leaders.
7. **Scalability:** Ability to elect leaders in large distributed systems.
8. **Complexity:** Analyzing time and message complexity of election algorithms.
9. **Use Cases:** Leader election is critical in distributed databases, cluster management, and resource allocation.
10. **Variations:** Different election algorithms designed for specific system requirements (e.g., crash vs. Byzantine failures).

10. **

Consensus Protocols**:

1. **Definition:** Algorithms that enable a group of distributed processes to agree on a common value.
2. **Paxos Algorithm:** A well-known consensus protocol for achieving agreement in distributed systems.
3. **Raft Algorithm:** An alternative to Paxos, designed for understandability and practicality.
4. **Quorum Systems:** A method to ensure agreement by requiring a subset of processes to respond.
5. **Failure Tolerance:** Mechanisms to achieve consensus despite process failures.
6. **Message Complexity:** Analyzing the number of messages required to reach consensus.
7. **Performance Metrics:** Evaluating latency and throughput in consensus protocols.
8. **Use Cases:** Critical for distributed databases, cloud computing, and fault-tolerant systems.
9. **Scalability Challenges:** Issues arising in large distributed systems regarding consensus.
10. **Variations:** Adaptations of consensus algorithms for different applications (e.g., asynchronous systems).